


# 601.418/618 (S24): Assignment 3A: Virtual Memory

 [jhuopsys.github.io/spring2024/assign/assign03a.html](https://jhuopsys.github.io/spring2024/assign/assign03a.html)

601.418/618 (S24): Assignment 3A: Virtual Memory

## Project 3: Virtual Memory - Part A

**Due:** Friday, April 12th by 11 pm

By now you should have some familiarity with the inner workings of Pintos. Your OS can properly handle multiple threads of execution with proper synchronization, and can load multiple user programs at once. However, the number and size of programs that can run is limited by the machine's main memory size. In this assignment, you will remove that limitation.

You will build this assignment also work with project 2 before you start work on problems in project 3.

You will continue to have previous assignment (

in project 2 should project 2 submission because the same

you did in the

## Background

### B.1 Source Files

You will work in the “v” “Makefile”s. The only setting “-DVM”. All code projects.

tains only “le” turns on the ed in earlier



You will probably be encountering just a few files for the first time:

“`devices/block.h`”

“`devices/block.c`”

Provides sector-based read and write access to block device. You will use this interface to access the swap partition as a block device.

### B.2 Memory Terminology


Careful definitions are needed to keep discussion of virtual memory from being confusing. Thus, we begin by presenting some terminology for memory and storage. Some of these terms should be familiar from project 2 (see section [4.1.4 Virtual Memory Layout](#)), but

much of it is new.

### B.2.1 Pages

---

A *page*, sometimes called a *virtual page*, is a continuous region of virtual memory 4,096 bytes (the *page size*) in length. A page must be *page-aligned*, that is, start on a virtual address evenly divisible by the page size. Thus, a 32-bit virtual address can be divided into a 20-bit *page number* and a 12-bit *page offset* (or just *offset*), like this:




Each process has an independent set of *user (virtual) pages*, which are those pages below virtual address `PHYS_BASE`, typically `0xc0000000` (3 GB). The set of *kernel (virtual) pages*, on the other hand, is global, remaining the same regardless of what thread or process is active. The kernel may access both user and kernel pages, but a user process may access only its own. See section [A.6 Virtual Addresses](#), for more information.

Pintos provides several functions for translating between virtual and physical addresses. See section [A.6 Virtual Addresses](#).

### B.2.2 Frames

---

A *frame*, sometimes called a *physical page*, is a continuous region of physical memory. Like a page, a frame must be *frame-aligned*. Thus, a 32-bit physical address can be divided into a 20-bit *frame number* and a 12-bit *frame offset* (or just *offset*), like this:



The 80x86 doesn't provide any way to directly access memory at a physical address. Pintos works around this by mapping kernel virtual memory directly to physical memory: the first page of kernel virtual memory is mapped to the first frame of physical memory, the second page to the second frame, and so on. Thus, frames can be accessed through kernel virtual memory.


Pintos provides functions for translating between physical addresses and kernel virtual addresses. See section [A.6 Virtual Addresses](#), for details.

### B.2.3 Page Tables

---

In Pintos, a *page table* is a data structure that the CPU uses to translate a virtual address to a physical address, that is, from a page to a frame. The page table format is dictated by the 80x86 architecture. Pintos provides page table management code in “[pagedir.c](#)” (see section [A.7 Page Table](#)).

The diagram below illustrates the relationship between pages and frames. The virtual address, on the left, consists of a page number and an offset. The page table translates the page number into a frame number, which is combined with the unmodified offset to obtain the physical address, on the right.



#### B.2.4 Swap Slots

A *swap slot* is a contiguous block of memory. Although hardware limits the number of swap slots and frames, swap slots are used to store pages so.

#### B.3 Resource Management

You will need to design a resource management system.

#### Supplemental page table

Enables page fault handling. See section [Managing the Supplemental Page Table](#).


#### Frame table

Allows efficient implementation of eviction policy. See section [B.5 Managing the Frame Table](#).

#### Swap table

Tracks usage of swap slots. See section [B.6 Managing the Swap Table](#).

You do not necessarily need to implement three completely distinct data structures: it may be convenient to wholly or partially merge related resources into a unified data structure.



partition.  
r than for pages  
wnside in doing

. See section [B.4](#)

To simplify your design, you may store these data structures in non-pageable memory. That means that you can be sure that pointers among them will remain valid.

Pintos includes a bitmap data structure in “`lib/kernel/bitmap.c`” and “`lib/kernel/bitmap.h`”. A bitmap is an array of bits, each of which can be true or false. Bitmaps are typically used to track usage in a set of (identical) resources: if resource *n* is in use, then bit *n* of the bitmap is true. Bitmaps are useful for many things, although you could extend their implementation to track other things.

Although more complex, they can offer significant benefits, they may also needlessly complicate your design. We recommend implementing any advanced features as part of your design.

The *supplemental page* is a page that contains all the data about each page. It is needed because the data is in a specific format. Such a data structure is often used to reduce confusion.

You may organize the supplemental page table as you wish. There are at least two basic approaches to its organization: in terms of segments or in terms of pages. Optionally, you may use the page table itself as an index to track the members of the supplemental page table. You will have to modify the Pintos page table implementation in “`pagedir.c`” to do so. We recommend this approach for advanced students only. See section [A.7.4.2 Page Table Entry Format](#), for more information.





The most important user of the supplemental page table is the page fault handler. In project 2, a page fault always indicated a bug in the kernel or a user program. In project 3, this is no longer true. Now, a page fault might only indicate that the page must be brought in from a file or swap. You will have to implement a more sophisticated page fault handler to handle these cases. Your page fault handler, which you should implement by modifying `page_fault()` in “`userprog/exception.c`”, needs to do roughly the following:

1. Locate the page that faulted in the supplemental page table. If the memory reference is valid, use the supplemental page table entry to locate the data that goes in the page, which might be in the file system, or in a swap slot, or it might simply be an all-zero page. If you implement sharing, the page's data might even already be in a page frame, but not in the page table.

If the supplemental page table indicates that the user process should not expect any data at the address it was trying to access, or if the page lies within kernel virtual memory, or if the access is an attempt to write to a read-only page, then the access is invalid. Any invalid access terminates the process and thereby frees all of its resources.

2. Obtain a frame to use for the page. See the [Frame Table](#), for details.

If you implement swapping, you must obtain a frame, in which case you must be careful to zero the page before using it.

3. Fetch the data into the page. If the data is in swap, zeroing it, etc.

If you implement sharing, you must obtain a frame, in which case no action is needed.

4. Point the page table entry to the physical page. You can use the function `page_table_entry_set_physical_page()`.



## B.5 Managing the Frame Table

The *frame table* contains one entry for each frame that contains a user page. Each entry in the frame table contains a pointer to the page, if any, that currently occupies it, and other data of your choice. The frame table allows Pintos to efficiently implement an eviction policy, by choosing a page to evict when no frames are free.

The frames used for user pages should be obtained from the "user pool," by calling `palloc_get_page(PAL_USER)`. You must use `PAL_USER` to avoid allocating from the "kernel pool," which could cause some test cases to fail unexpectedly (see [Why PAL\\_USER?](#)). If you modify “`palloc.c`” as part of your frame table implementation, be sure to retain the distinction between the two pools.

The most important operation on the frame table is obtaining an unused frame. This is easy when a frame is free. When none is free, a frame must be made free by evicting some page from its frame.

If no frame can be evicted without allocating a swap slot, but swap is full, panic the kernel. Real OSes apply a wide range of policies to recover from or prevent such situations, but these policies are beyond the scope of this project.

The process of eviction comprises roughly the following steps:

1. Choose a frame to evict, using your page replacement algorithm. The "accessed" and "dirty" bits in the page table, described below, will come in handy.
2. Remove references to the frame from any page table that refers to it.

Unless you have implemented sharing, only a single page should refer to a frame at any given time.


3. If necessary, write the page to disk.

The evicted frame may be reused.

### B.5.1 Accessed and Dirty Bits

80x86 hardware provides a way to set the accessed and dirty bits in the page table. On any read or write to a page, the CPU sets the accessed bit. On any write, the CPU sets the dirty bit. The OS may do so.

You need to be aware of these bits when you manage the frame table. When an alias is created, only one page table entry is needed. The accessed and dirty bits for the other aliases are updated in the frame table.



placement page. On any PTE, and on any to 0, but the OS

to the same are updated in accessed and dirty

In Pintos, every user virtual page is aliased to its kernel virtual page. You must manage these aliases somehow. For example, your code could check and update the accessed and dirty bits for both addresses. Alternatively, the kernel could avoid the problem by only accessing user data through the user virtual address.

Other aliases should only arise if you implement sharing for extra credit (see [VM Extra Credit](#)), or if there is a bug in your code.

See section [A.7.3 Accessed and Dirty Bits](#), for details of the functions to work with accessed and dirty bits.

## B.6 Managing the Swap Table